

The Application Layer – HTTP and FTP

File Transfer Protocol (FTP)

- Allows a user to copy files to/from remote hosts
 - Client program connects to FTP server
 - ... provides a login id and password
 - ... allows the user to explore the directories
 - ... and download and upload files with the server
- A predecessor of the Web (RFC 959 in 1985)
 - Requires user to know the name of the server machine
 - ... and have an account on the machine
 - ... and find the directory where the files are stored
 - ... and know whether the file is text or binary
 - ... and know what tool to run to render and edit the file
- That is, no URL, hypertext, and helper applications

Why Study FTP?


- Helps cement familiarity with text/status-code based protocols
- Illustrates use of **multiple concurrent connections**
 - One for control (commands & replies)
 - Depending on command, can be additional one for data
- Illustrates **reversal of roles**
 - For data connection, FTP user's process can play the server role, FTP server can play the client role

Example commands

- Authentication
 - USER: specify the user name to log in as
 - PASS: specify the user's password
- Exploring the files
 - LIST: list the files for the given file specification
 - CWD: change to the given directory
- Downloading and uploading files
 - TYPE: set type to ASCII (A) or binary image (I)
 - RETR: retrieve the given file
 - STOR: upload the given file
- Closing the connection
 - QUIT: close the FTP connection


FTP Data Transfer

- Separate data connection
 - To send lists of files (LIST)
 - To retrieve a file (RETR)
 - To upload a file (STOR)



Creating the Data Connection


- Client acts like a server
 - Creates a socket
 - Assigned an ephemeral port number by the kernel
 - Listens on socket
 - Waits to hear from FTP server



Creating Data Connection (cont.)

- But, the server doesn't know the port number
 - So after starting to listen, client tells it to the server
 - Using the PORT command on the control connection
 - Server can tell the client a port to connect to using PASV or EPSV

PORT <IP address, port #>



Why Out-of-Band Control?

- Avoids need to mark the end of the data transfer
 - Data transfer ends by closing of data connection
 - Yet, the control connection stays up
- Aborting a data transfer
 - Can abort a transfer without killing the control connection
 - ... which avoids requiring the user to log in again
 - Done with an ABOR on the control connection
- Third-party file transfer between two hosts
 - Data connection could go to a different host
 - ... by sending a different client IP address to the server
 - e.g., a user can coordinate a transfer between two servers
 - But: this is rarely needed, and presents security issues

HTTP

- Server listens on a port (by default, 80)
- On connection, waits for a request
- Protocol (but not data) is in ASCII
- Sends response, maybe closes connection (client can ask it to stay open)

Parsing a URL

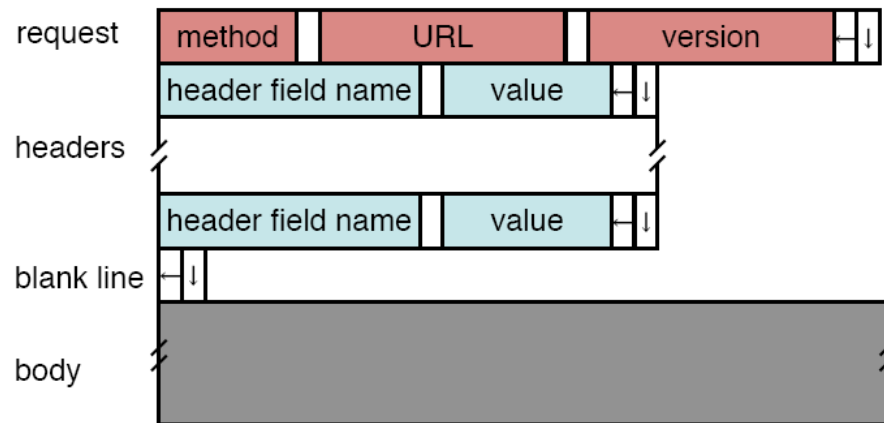
http://www.niit.edu.pk/~tahir/tcpip/index.html

Protocol Host File path on host



The diagram illustrates the parsing of the URL http://www.niit.edu.pk/~tahir/tcpip/index.html. The URL is divided into three main components, each highlighted with a different color and underlined. The first component, http, is red and labeled 'Protocol'. The second component, www.niit.edu.pk, is teal and labeled 'Host'. The third component, ~tahir/tcpip/index.html, is purple and labeled 'File path on host'. Arrows point from the labels below to their respective components in the URL.

HTTP Request Format



- Request types: GET, PUT, POST, HEAD, DELETE
- A trivial browser request: `http://localhost:8000`

Other useful header fields

- Range: Request a partial range of data
- Authorization: Present authorization credentials to a server (not HTTPS)
- Proxy-Authorization: Present proxy credentials to a proxy server
- Referer: URL of the web page the user was on, when the HTTP request was made

HTTP Performance


- What matters most?
- Different kinds of requests
 - Lots of small requests (loading a web page)
 - Big request (fetching a download)
- Require different solutions

Small requests

- Latency matters
- Governed by RTT between hosts
- Two major causes of delay:
 - Opening a TCP connection
 - Data response-request
- Solutions:
 - Persistent connections
 - Pre-fetching
 - Others??

Big requests

- Problem is throughput on bottleneck links (usually edge links)
- Use an HTTP proxy cache or mirror
 - Can also improve latency!



Client-server vs Peer-to-peer

- Server can be a bottleneck
 - Download time can scale down $O(n)$ with n clients
 - Scaling up server bandwidth can be expensive (CDNs)
 - Slashdotting/flash crowds
- Peer-to-peer: get a bunch of end-hosts to collaboratively distribute content
- A common peer-to-peer challenge is finding whom to collaborate with